

Rain Prediction using ML

Time Line of the Project :

- Data Analysis
- Handling Missing Values
- Handling Categorical Variables
- Feature Engineering
- Model Building using ML
- Model Building using Auto ML i.e PyCaret

Importing Libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import
accuracy_score, confusion_matrix, classification_report
from sklearn import metrics
import math
%matplotlib inline
```

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

```
df= pd.read_csv("/content/weatherAUS.csv")
pd.set_option("display.max_columns", None)
```

```
df
```

```
{"type": "dataframe", "variable_name": "df"}
```

```
df.nunique()
```

Date	3436
Location	49
MinTemp	389
MaxTemp	505
Rainfall	681
Evaporation	358
Sunshine	145
WindGustDir	16
WindGustSpeed	67

WindDir9am	16
WindDir3pm	16
WindSpeed9am	43
WindSpeed3pm	44
Humidity9am	101
Humidity3pm	101
Pressure9am	546
Pressure3pm	549
Cloud9am	10
Cloud3pm	10
Temp9am	441
Temp3pm	502
RainToday	2
RainTomorrow	2

dtype: int64

Discrete Variable are countable in finite amount of time while numerical variable are too much in number to count

```
# Numerical features
num_var = df.select_dtypes(include=['int64',
'float64']).columns.tolist()

# Discrete numerical features (few unique values)
discrete_var = [col for col in num_var if df[col].nunique() <= 25]

# Continuous numerical features
cont_var = [col for col in num_var if col not in discrete_var]

# Categorical features (object or category type)
categ_var = df.select_dtypes(include=['object']).columns.tolist()

df[categ_var]

{"type": "dataframe"}
```

Handling Missing Values

```
df.isnull().sum()

Date      0
Location  0
MinTemp   1485
MaxTemp   1261
Rainfall  3261
Evaporation  62790
Sunshine  69835
WindGustDir  10326
WindGustSpeed  10263
WindDir9am  10566
```

WindDir3pm	4228
WindSpeed9am	1767
WindSpeed3pm	3062
Humidity9am	2654
Humidity3pm	4507
Pressure9am	15065
Pressure3pm	15028
Cloud9am	55888
Cloud3pm	59358
Temp9am	1767
Temp3pm	3609
RainToday	3261
RainTomorrow	3267

dtype: int64

the percentage of missing values in each column.

```
df.isnull().sum()*100/len(df)
```

Date	0.000000
Location	0.000000
MinTemp	1.020899
MaxTemp	0.866905
Rainfall	2.241853
Evaporation	43.166506
Sunshine	48.009762
WindGustDir	7.098859
WindGustSpeed	7.055548
WindDir9am	7.263853
WindDir3pm	2.906641
WindSpeed9am	1.214767
WindSpeed3pm	2.105046
Humidity9am	1.824557
Humidity3pm	3.098446
Pressure9am	10.356799
Pressure3pm	10.331363
Cloud9am	38.421559
Cloud3pm	40.807095
Temp9am	1.214767
Temp3pm	2.481094
RainToday	2.241853
RainTomorrow	2.245978

dtype: float64

```
def find_var_type(var):
```

```
    if var in discrete_var:
        print("{} is a Numerical Variable, Discrete in
nature".format(var))
```

```

elif var in cont_var :
    print("{} is a Numerical Variable, Continuous in
nature".format(var))
else :
    print("{} is a Categorical Variable".format(var))

find_var_type('Cloud3pm')

Cloud3pm is a Numerical Variable, Discrete in nature

```

Random Sample Imputation for the our variables which are having the most percentage of Null Values

```

def RandomSampleImputation(df, feature):
    # Randomly sample from non-null values
    random_sample = df[feature].dropna().sample(
        df[feature].isnull().sum(), random_state=0, replace=True
    )
    # Align the sampled index to the null index
    random_sample.index = df[df[feature].isnull()].index

    # Fill in the missing values with the sampled values
    df.loc[df[feature].isnull(), feature] = random_sample

```

```

RandomSampleImputation(df, "Cloud9am")
RandomSampleImputation(df, "Cloud3pm")
RandomSampleImputation(df, "Evaporation")
RandomSampleImputation(df, "Sunshine")

```

```
df.isnull().sum()*100/len(df)
```

Date	0.000000
Location	0.000000
MinTemp	1.020899
MaxTemp	0.866905
Rainfall	2.241853
Evaporation	0.000000
Sunshine	0.000000
WindGustDir	7.098859
WindGustSpeed	7.055548
WindDir9am	7.263853
WindDir3pm	2.906641
WindSpeed9am	1.214767
WindSpeed3pm	2.105046
Humidity9am	1.824557
Humidity3pm	3.098446
Pressure9am	10.356799
Pressure3pm	10.331363
Cloud9am	0.000000
Cloud3pm	0.000000

```
Temp9am      1.214767
Temp3pm      2.481094
RainToday    2.241853
RainTomorrow 2.245978
dtype: float64
```

```
find_var_type('RainToday')
```

```
RainToday is a Categorical Variable
```

replace the null values of all the continuous feature which are having less number of null values

```
def MeanImputation(df, feature):
    df[feature]= df[feature]
    mean= df[feature].mean()
    df[feature]= df[feature].fillna(mean)
```

```
MeanImputation(df, 'Pressure3pm')
```

```
MeanImputation(df, 'Pressure9am')
```

```
MeanImputation(df, 'MinTemp')
```

```
MeanImputation(df, 'MaxTemp')
```

```
MeanImputation(df, 'Rainfall')
```

```
MeanImputation(df, 'WindGustSpeed')
```

```
MeanImputation(df, 'WindSpeed9am')
```

```
MeanImputation(df, 'WindSpeed3pm')
```

```
MeanImputation(df, 'Pressure9am')
```

```
MeanImputation(df, 'Humidity9am')
```

```
MeanImputation(df, 'Humidity3pm')
```

```
MeanImputation(df, 'Temp3pm')
```

```
MeanImputation(df, 'Temp9am')
```

```
df.isnull().sum()*100/len(df)
```

```
Date      0.000000
Location  0.000000
MinTemp    0.000000
MaxTemp    0.000000
Rainfall   0.000000
Evaporation 0.000000
Sunshine   0.000000
WindGustDir 7.098859
WindGustSpeed 0.000000
WindDir9am 7.263853
WindDir3pm 2.906641
WindSpeed9am 0.000000
WindSpeed3pm 0.000000
Humidity9am 0.000000
Humidity3pm 0.000000
```

```
Pressure9am      0.000000
Pressure3pm      0.000000
Cloud9am         0.000000
Cloud3pm         0.000000
Temp9am          0.000000
Temp3pm          0.000000
RainToday        2.241853
RainTomorrow     2.245978
dtype: float64
```

Plotting a HeatMap for the numerical values

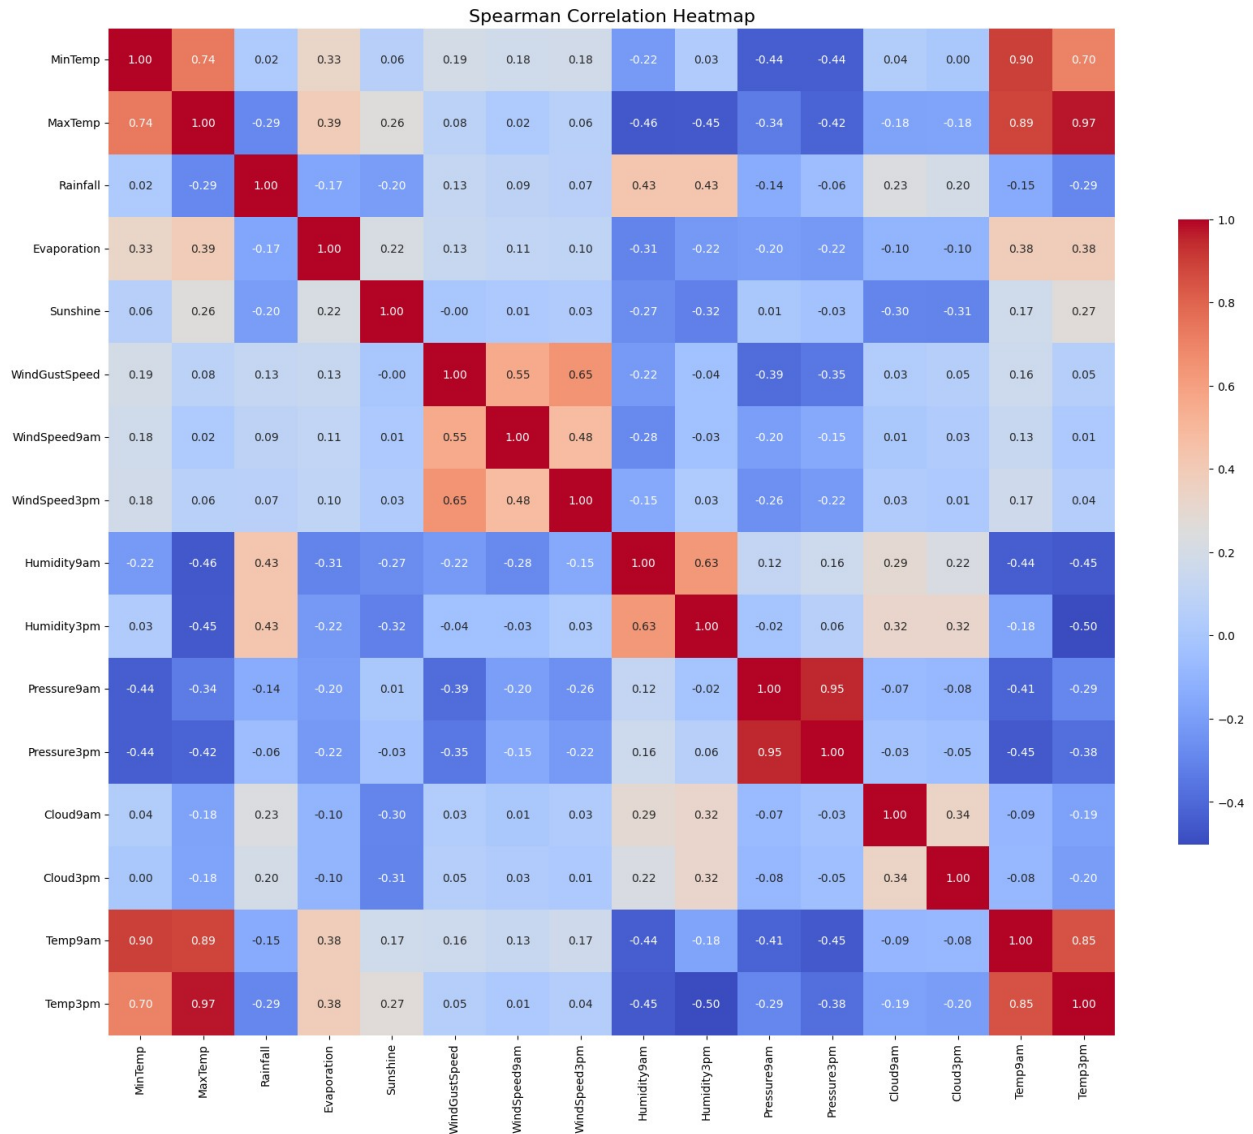
```
import matplotlib.pyplot as plt
import seaborn as sns

# Select only numeric columns for correlation
numeric_df = df.select_dtypes(include=['number'])

# Compute Spearman correlation matrix
corrmat = numeric_df.corr(method="spearman")

# Set plot size
plt.figure(figsize=(20, 20))

# Plot heatmap
sns.heatmap(corrmat, annot=True, fmt=".2f", cmap="coolwarm",
            square=True, cbar_kws={"shrink": 0.5})
plt.title("Spearman Correlation Heatmap", fontsize=16)
plt.show()
```



Analysis for Continuous variables

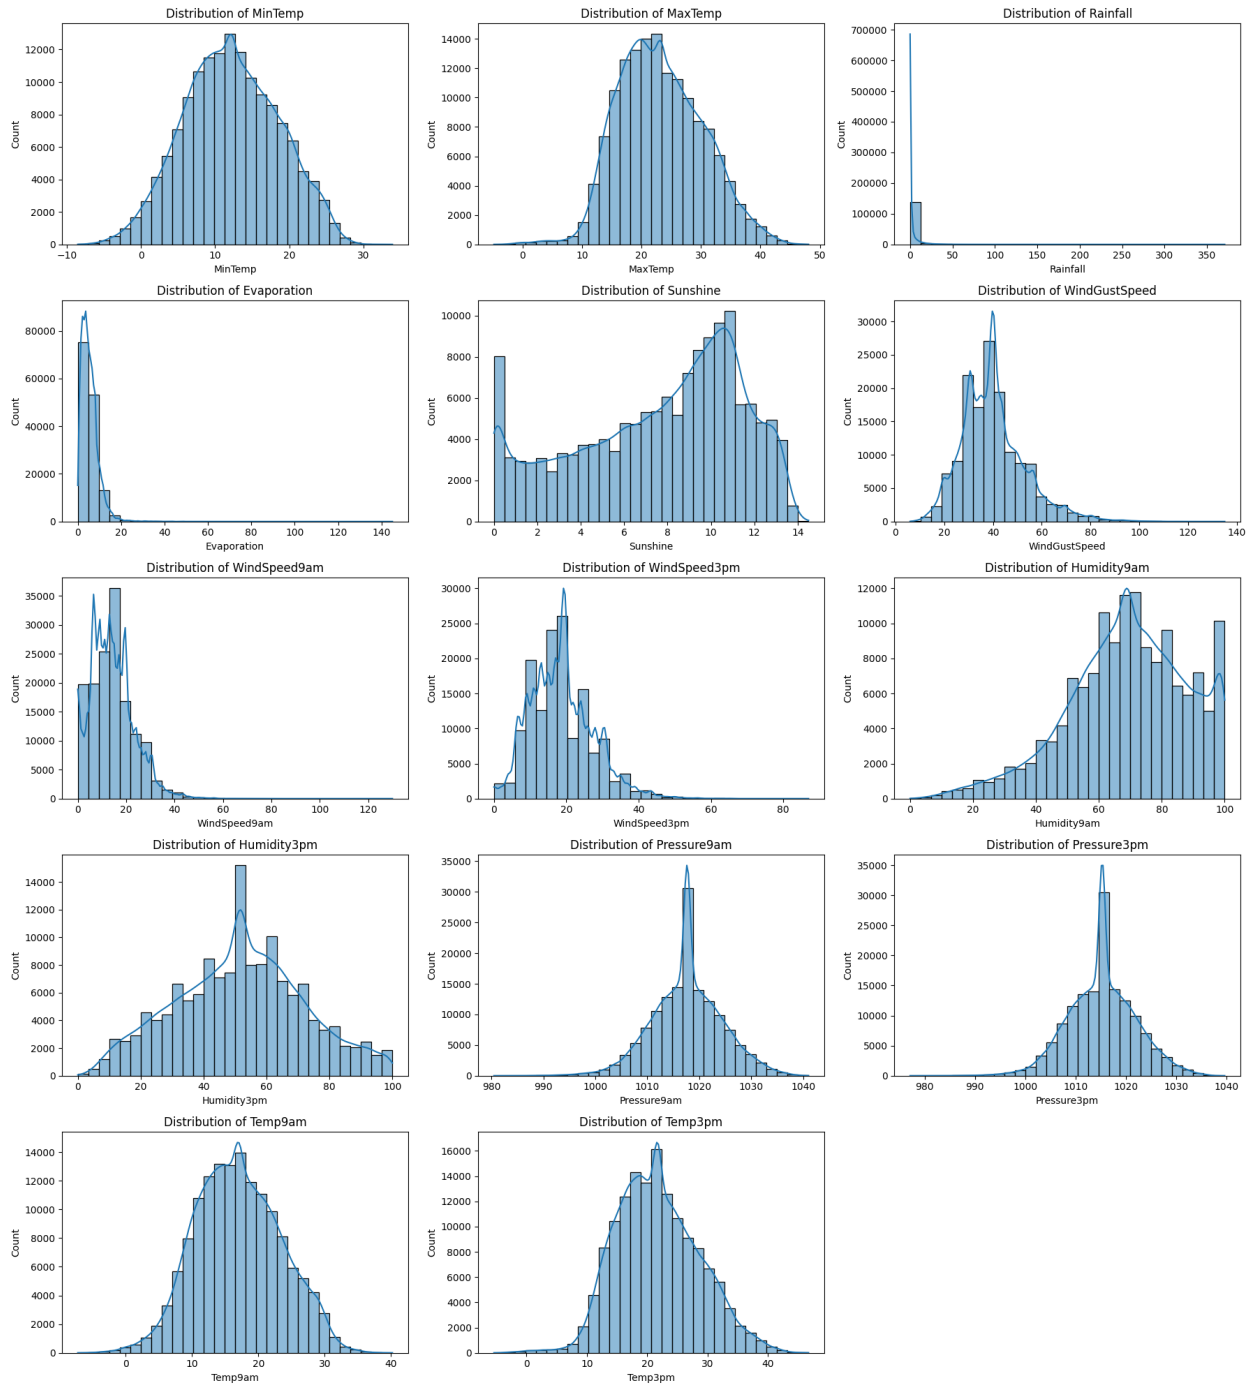
```
# Number of plots
n = len(cont_var)
cols = 3 # Number of columns in the grid
rows = math.ceil(n / cols)

# Set figure size
plt.figure(figsize=(cols * 6, rows * 4))

for idx, feature in enumerate(cont_var):
    plt.subplot(rows, cols, idx + 1)
    sns.histplot(data=df, x=feature, kde=True, bins=30)
    plt.xlabel(feature)
```

```
plt.ylabel("Count")
plt.title(f"Distribution of {feature}")
```

```
plt.tight_layout()
plt.show()
```

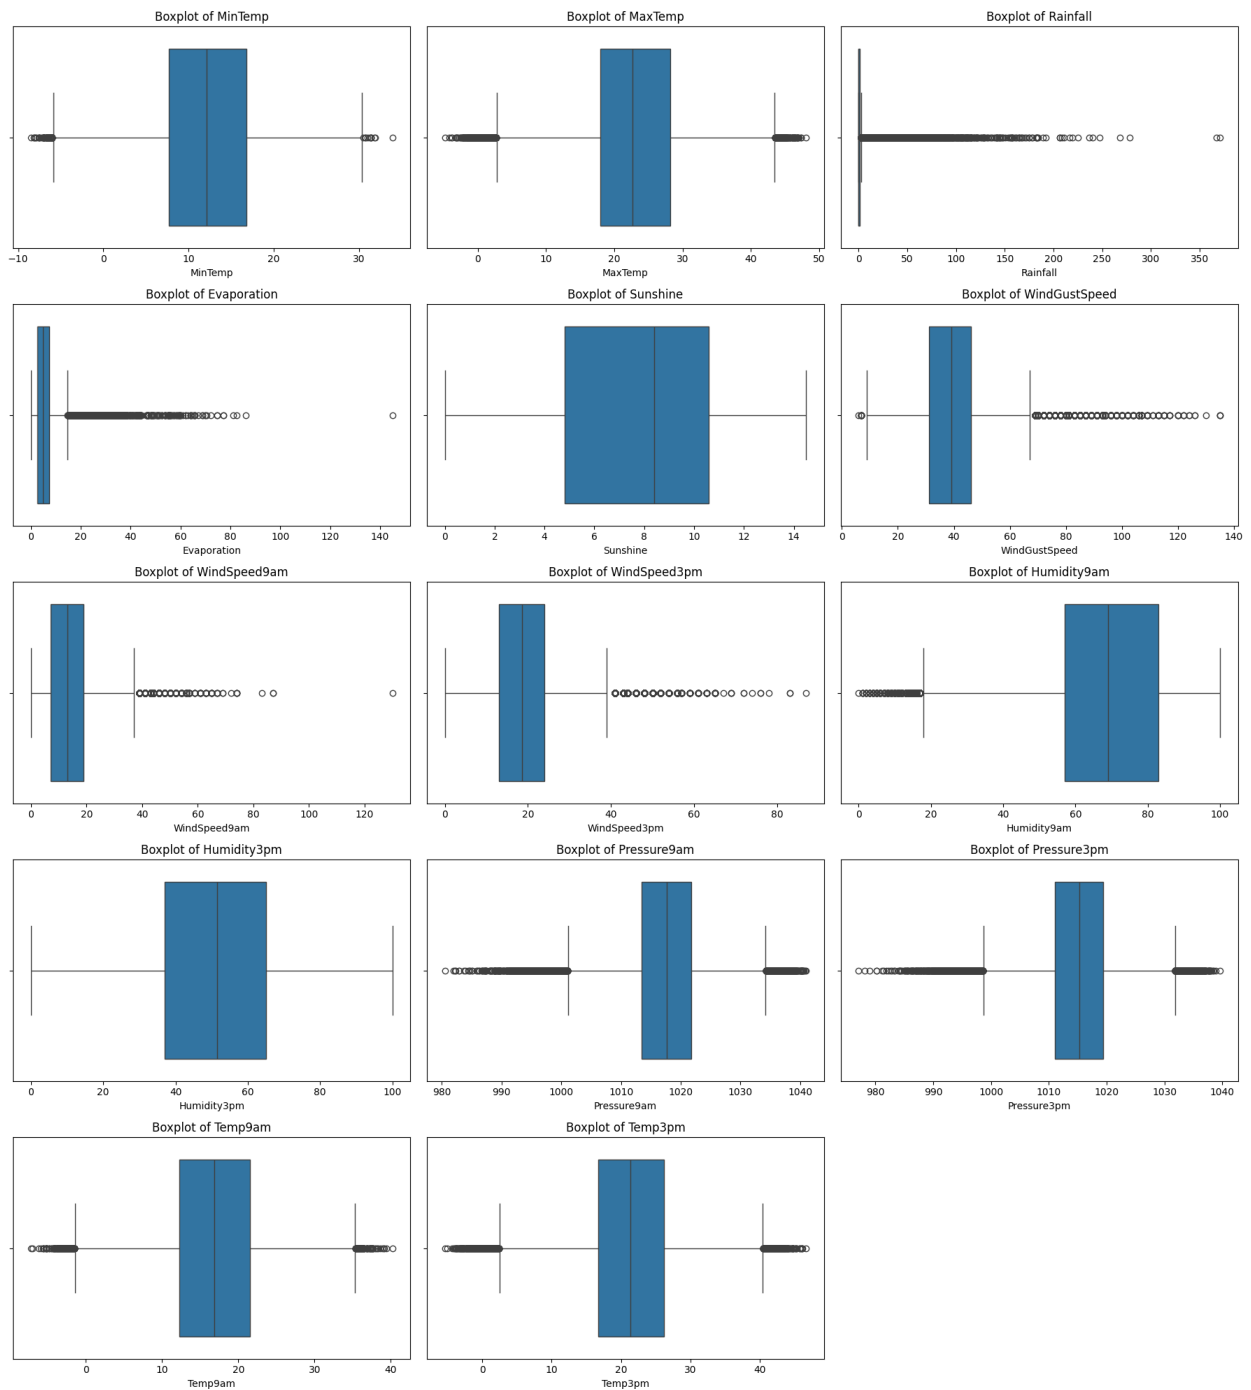



```
# Number of plots
n = len(cont_var)
cols = 3 # Number of columns in the grid
rows = math.ceil(n / cols)

# Set figure size
plt.figure(figsize=(cols * 6, rows * 4))

for idx, feature in enumerate(cont_var):
    plt.subplot(rows, cols, idx + 1)
    sns.boxplot(x=df[feature])
    plt.title(f"Boxplot of {feature}")
    plt.xlabel(feature)

plt.tight_layout()
plt.show()
```



One Hot Encoding

```
df["RainToday"] = pd.get_dummies(df["RainToday"], drop_first = True)
df["RainTomorrow"] = pd.get_dummies(df["RainTomorrow"], drop_first = True)
df
{"type": "dataframe", "variable_name": "df"}
```

Lable Encoding

```
for feature in categ_var:
    print(feature, (df.groupby([feature])
["RainTomorrow"].mean().sort_values(ascending = False)).index)

Date Index(['2007-12-15', '2007-11-01', '2007-11-02', '2007-12-27',
'2008-01-18',
'2008-02-12', '2008-02-07', '2007-12-24', '2008-01-19', '2008-
01-12',
...
'2009-10-21', '2009-10-22', '2007-12-14', '2007-12-17', '2007-
12-18',
'2007-12-23', '2007-11-07', '2007-11-05', '2007-11-12', '2007-
11-06'],
dtype='object', name='Date', length=3436)
Location Index(['Portland', 'Walpole', 'Cairns', 'Dartmoor',
'NorfolkIsland',
'MountGambier', 'Albany', 'Witchcliffe', 'CoffsHarbour',
'MountGinini',
'NorahHead', 'Darwin', 'Sydney', 'SydneyAirport', 'Ballarat',
'GoldCoast', 'Watsonia', 'Newcastle', 'Hobart', 'Wollongong',
'Williamtown', 'Launceston', 'Brisbane', 'MelbourneAirport',
'Adelaide',
'Sale', 'Albury', 'Perth', 'Melbourne', 'Nuriootpa', 'Penrith',
'BadgerysCreek', 'PerthAirport', 'Tuggeranong', 'Richmond',
'Bendigo',
'Canberra', 'WaggaWagga', 'Townsville', 'Katherine',
'PearceRAAF',
'SalmonGums', 'Nhil', 'Moree', 'Cobar', 'Mildura',
'AliceSprings',
'Uluru', 'Woomera'],
dtype='object', name='Location')
WindGustDir Index(['NNW', 'NW', 'WNW', 'N', 'W', 'WSW', 'NNE', 'S',
'SSW', 'SW', 'SSE',
'NE', 'SE', 'ESE', 'ENE', 'E'],
dtype='object', name='WindGustDir')
WindDir9am Index(['NNW', 'N', 'NW', 'NNE', 'WNW', 'W', 'WSW', 'SW',
'SSW', 'NE', 'S',
'SSE', 'ENE', 'SE', 'ESE', 'E'],
dtype='object', name='WindDir9am')
WindDir3pm Index(['NW', 'NNW', 'N', 'WNW', 'W', 'NNE', 'WSW', 'SSW',
'S', 'SW', 'SE',
'NE', 'SSE', 'ENE', 'E', 'ESE'],
dtype='object', name='WindDir3pm')
RainToday Index([True, False], dtype='bool', name='RainToday')
RainTomorrow Index([True, False], dtype='bool', name='RainTomorrow')

windgustdir = {'NNW':0, 'NW':1, 'WNW':2, 'N':3, 'W':4, 'WSW':5,
'NNE':6, 'S':7, 'SSW':8, 'SW':9, 'SSE':10,
```

```

        'NE':11, 'SE':12, 'ESE':13, 'ENE':14, 'E':15}
winddir9am = {'NNW':0, 'N':1, 'NW':2, 'NNE':3, 'WNW':4, 'W':5,
              'WSW':6, 'SW':7, 'SSW':8, 'NE':9, 'S':10,
              'SSE':11, 'ENE':12, 'SE':13, 'ESE':14, 'E':15}
winddir3pm = {'NW':0, 'NNW':1, 'N':2, 'WNW':3, 'W':4, 'NNE':5,
              'WSW':6, 'SSW':7, 'S':8, 'SW':9, 'SE':10,
              'NE':11, 'SSE':12, 'ENE':13, 'E':14, 'ESE':15}
df["WindGustDir"] = df["WindGustDir"].map(windgustdir)
df["WindDir9am"] = df["WindDir9am"].map(winddir9am)
df["WindDir3pm"] = df["WindDir3pm"].map(winddir3pm)

df["WindGustDir"] =
df["WindGustDir"].fillna(df["WindGustDir"].value_counts().index[0])
df["WindDir9am"] =
df["WindDir9am"].fillna(df["WindDir9am"].value_counts().index[0])
df["WindDir3pm"] =
df["WindDir3pm"].fillna(df["WindDir3pm"].value_counts().index[0])

df.isnull().sum()*100/len(df)

Date                0.0
Location            0.0
MinTemp             0.0
MaxTemp             0.0
Rainfall            0.0
Evaporation         0.0
Sunshine            0.0
WindGustDir         0.0
WindGustSpeed       0.0
WindDir9am          0.0
WindDir3pm          0.0
WindSpeed9am        0.0
WindSpeed3pm        0.0
Humidity9am          0.0
Humidity3pm          0.0
Pressure9am          0.0
Pressure3pm          0.0
Cloud9am             0.0
Cloud3pm             0.0
Temp9am              0.0
Temp3pm              0.0
RainToday            0.0
RainTomorrow         0.0
dtype: float64

df.head()

{"type": "dataframe", "variable_name": "df"}

```

We have removed all the null values and handled with categorical data

Now we will do the Label Encoding for our Location according to our Target variable

```
df_loc = df.groupby(["Location"])
["RainTomorrow"].value_counts().sort_values().unstack()

df_loc.head()

{"summary": "{\n  \"name\": \"df_loc\",\n  \"rows\": 49,\n  \"fields\": [\n    {\n      \"column\": \"Location\",\n      \"properties\": {\n        \"dtype\": \"string\",\n        \"num_unique_values\": 49,\n        \"samples\": [\n          \"Darwin\",\n          \"Williamtown\",\n          \"Wollongong\"\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": false,\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 315,\n        \"min\": 1313,\n        \"max\": 2807,\n        \"num_unique_values\": 46,\n        \"samples\": [\n          1462,\n          2196,\n          2417\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      },\n      \"column\": true,\n      \"properties\": {\n        \"dtype\": \"number\",\n        \"std\": 219,\n        \"min\": 116,\n        \"max\": 1095,\n        \"num_unique_values\": 49,\n        \"samples\": [\n          852,\n          700,\n          713\n        ],\n        \"semantic_type\": \"\",\n        \"description\": \"\"\n      }\n    ]\n  },\n  \"type\": \"dataframe\",\n  \"variable_name\": \"df_loc\"}
```

```
import pandas as pd

row = df.iloc[1] # This is a Series
numeric_row = pd.to_numeric(row, errors='coerce') # Convert values to numbers; non-numeric become NaN
numeric_row = numeric_row.dropna() # Drop NaNs
numeric_row.sort_values(ascending=False)
```

Pressure9am	1010.6
Pressure3pm	1007.8
WindGustSpeed	44.0
Humidity9am	44.0
MaxTemp	25.1
Humidity3pm	25.0
Temp3pm	24.3
WindSpeed3pm	22.0
Temp9am	17.2
Sunshine	11.9
Cloud3pm	8.0
MinTemp	7.4
WindDir3pm	6.0

WindSpeed9am	4.0
Evaporation	2.6
Cloud9am	2.0
WindGustDir	2.0
Rainfall	0.0
WindDir9am	0.0
RainToday	0.0
RainTomorrow	0.0

Name: 1, dtype: float64

```
# Assuming df_loc is grouped by "Location" and then counts of
"RainTomorrow" are calculated
df_loc = df.groupby(["Location"])
["RainTomorrow"].value_counts().sort_values().unstack()

# # Sorting the counts for the specific location (in this case, 1
refers to the second location, change if needed)
# sorted_values = df_loc.iloc[1].sort_values(ascending=False)

# # Display the sorted values
# print(sorted_values)

df_loc.head()

{"summary":{"\n  \"name\": \"df_loc\", \n  \"rows\": 49, \n  \"fields\":
[\n    {\n      \"column\": \"Location\", \n      \"properties\": {\n
\"dtype\": \"string\", \n      \"num_unique_values\": 49, \n
\"samples\": [\n        \"Darwin\", \n        \"Williamtown\", \n
\"Wollongong\", \n      ], \n      \"semantic_type\": \"\", \n
\"description\": \"\" \n    }, \n    {\n      \"column\":
false, \n      \"properties\": {\n        \"dtype\": \"number\", \n
\"std\": 315, \n        \"min\": 1313, \n        \"max\": 2807, \n
\"num_unique_values\": 46, \n        \"samples\": [\n          1462, \n
2196, \n          2417, \n        ], \n        \"semantic_type\": \"\", \n
\"description\": \"\" \n      }, \n      {\n        \"column\":
true, \n        \"properties\": {\n          \"dtype\": \"number\", \n
\"std\": 219, \n          \"min\": 116, \n          \"max\": 1095, \n
\"num_unique_values\": 49, \n          \"samples\": [\n            852, \n
700, \n            713, \n          ], \n          \"semantic_type\": \"\", \n
\"description\": \"\" \n        } \n      ] \n    } \n  ],
\"type\":\"dataframe\", \"variable_name\":\"df_loc\"}
```

```
df_loc[True].sort_values(ascending=False)
```

Location	
Portland	1095
Cairns	950
Walpole	949
Dartmoor	922
MountGambier	920
NorfolkIsland	919

Albany	902
Witchcliffe	879
CoffsHarbour	869
Sydney	865
Darwin	852
MountGinini	819
NorahHead	808
Ballarat	781
GoldCoast	775
SydneyAirport	774
Hobart	761
Watsonia	738
Newcastle	731
Wollongong	713
Brisbane	709
Williamstown	700
Launceston	699
Adelaide	688
MelbourneAirport	653
Perth	645
Sale	643
Melbourne	636
Canberra	629
Albury	618
Penrith	595
Nuriootpa	592
BadgerysCreek	583
Tuggeranong	568
PerthAirport	567
Bendigo	562
Richmond	560
WaggaWagga	536
Townsville	519
PearceRAAF	505
SalmonGums	472
Moree	394
Cobar	386
Mildura	327
Katherine	265
AliceSprings	244
Nhil	242
Woomera	202
Uluru	116

Name: True, dtype: int64

```
df_loc[True].sort_values(ascending=False).index
```

```
Index(['Portland', 'Cairns', 'Walpole', 'Dartmoor', 'MountGambier',
      'NorfolkIsland', 'Albany', 'Witchcliffe', 'CoffsHarbour',
      'Sydney',
```

```

        'Darwin', 'MountGinini', 'NorahHead', 'Ballarat', 'GoldCoast',
        'SydneyAirport', 'Hobart', 'Watsonia', 'Newcastle',
        'Wollongong',
        'Brisbane', 'Williamstown', 'Launceston', 'Adelaide',
        'MelbourneAirport',
        'Perth', 'Sale', 'Melbourne', 'Canberra', 'Albury', 'Penrith',
        'Nuriootpa', 'BadgerysCreek', 'Tuggeranong', 'PerthAirport',
        'Bendigo',
        'Richmond', 'WaggaWagga', 'Townsville', 'PearceRAAF',
        'SalmonGums',
        'Moree', 'Cobar', 'Mildura', 'Katherine', 'AliceSprings',
        'Nhil',
        'Woomera', 'Uluru'],
        dtype='object', name='Location')

```

```
len(df_loc[True].sort_values(ascending=False))
```

```
49
```

```

mapped_location = {'Portland':1, 'Cairns':2, 'Walpole':3,
'Dartmoor':4, 'MountGambier':5,
        'NorfolkIsland':6, 'Albany':7, 'Witchcliffe':8,
'CoffsHarbour':9, 'Sydney':10,
        'Darwin':11, 'MountGinini':12, 'NorahHead':13, 'Ballarat':14,
'GoldCoast':15,
        'SydneyAirport':16, 'Hobart':17, 'Watsonia':18, 'Newcastle':19,
'Wollongong':20,
        'Brisbane':21, 'Williamstown':22, 'Launceston':23,
'Adelaide':24, 'MelbourneAirport':25,
        'Perth':26, 'Sale':27, 'Melbourne':28, 'Canberra':29,
'Albury':30, 'Penrith':31,
        'Nuriootpa':32, 'BadgerysCreek':33, 'Tuggeranong':34,
'PerthAirport':35, 'Bendigo':36,
        'Richmond':37, 'WaggaWagga':38, 'Townsville':39,
'PearceRAAF':40, 'SalmonGums':41,
        'Moree':42, 'Cobar':43, 'Mildura':44, 'Katherine':45,
'AliceSprings':46, 'Nhil':47,
        'Woomera':48, 'Uluru':49}
df["Location"] = df["Location"].map(mapped_location)

```

Mapping Data

```

# df["Date"] = pd.to_datetime(df["Date"], format = "%Y-%m-%dT", errors
= "coerce")
# df["Date_month"] = df["Date"].dt.month
# df["Date_day"] = df["Date"].dt.day

df["Date"] = pd.to_datetime(df["Date"], format="%Y-%m-%d",
errors="coerce")

```

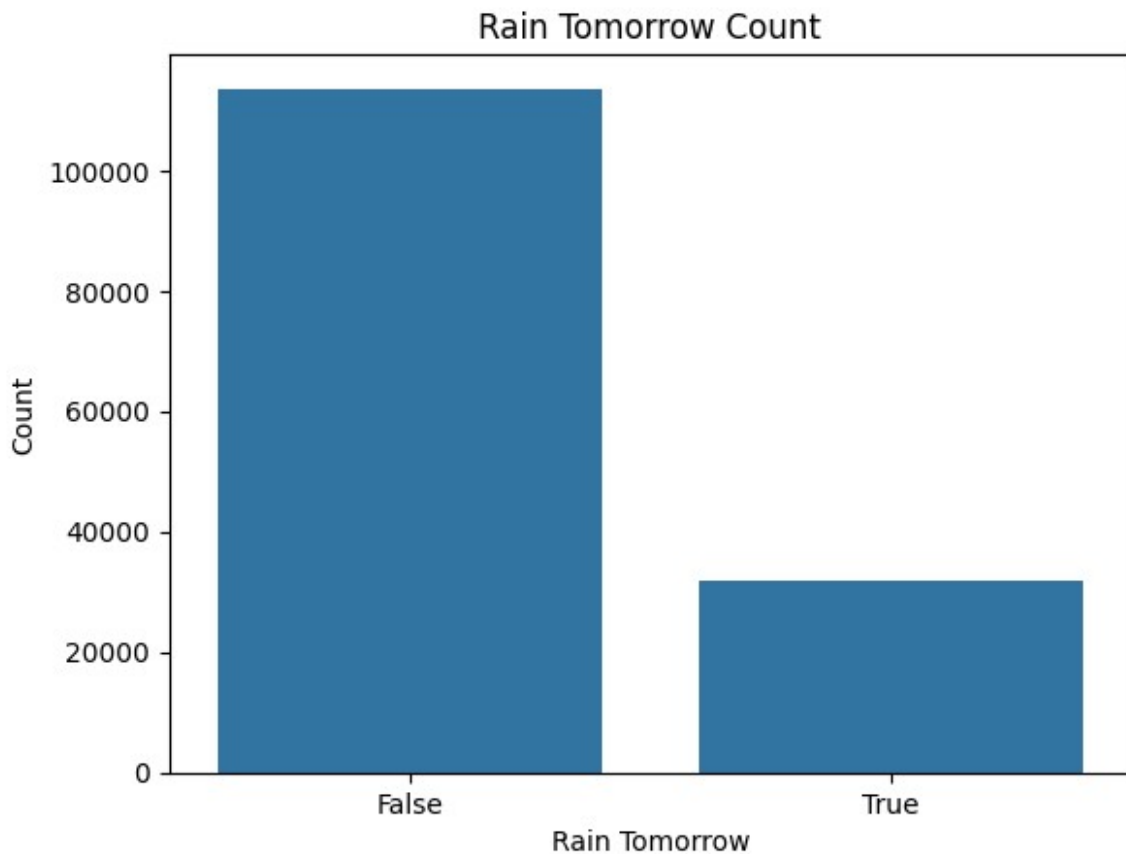


```
df["Date_month"] = df["Date"].dt.month
df["Date_day"] = df["Date"].dt.day

df.head()

{"type": "dataframe", "variable_name": "df"}

sns.countplot(x="RainTomorrow", data=df)
plt.title("Rain Tomorrow Count")
plt.xlabel("Rain Tomorrow")
plt.ylabel("Count")
plt.show()
```



```
df= df.drop(['Date'],axis=1)
df.head()

{"type": "dataframe", "variable_name": "df"}
```

Plotting Q-Q Plot

```
import scipy.stats as stats
import pylab
```

```

def plot_curve_grid(df, cont_var, rows=5, cols=4):
    fig, axes = plt.subplots(rows * 2, cols, figsize=(cols * 5, rows *
4))
    axes = axes.flatten()

    for i, feature in enumerate(cont_var):
        # Histogram
        axes[2 * i].hist(df[feature].dropna(), bins=30,
color='skyblue')
        axes[2 * i].set_title(f'{feature} - Histogram')

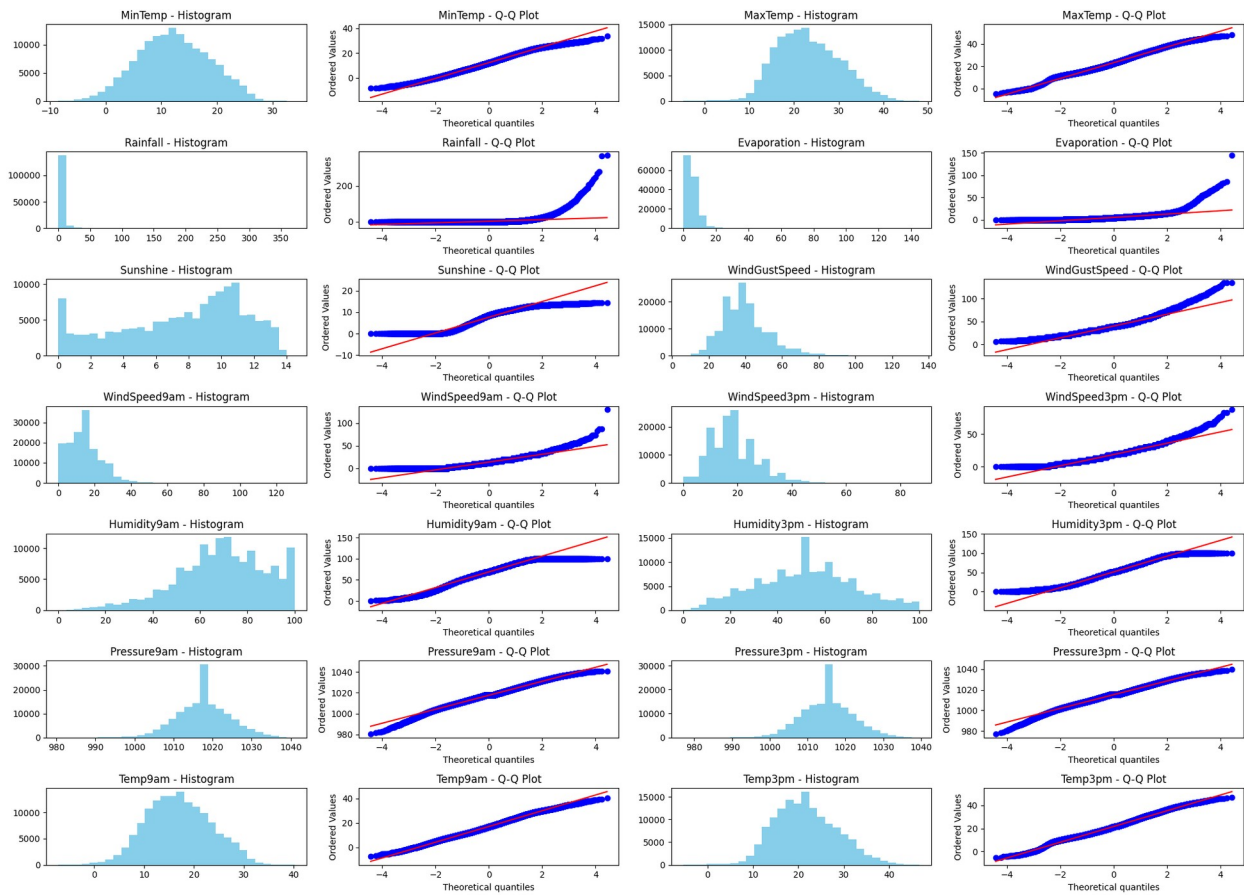
        # Q-Q plot
        stats.probplot(df[feature].dropna(), dist="norm", plot=axes[2
* i + 1])
        axes[2 * i + 1].set_title(f'{feature} - Q-Q Plot')

        # Hide any unused subplots
        for j in range(2 * len(cont_var), len(axes)):
            axes[j].axis('off')

    plt.tight_layout()
    plt.show()

# Call the function
plot_curve_grid(df, cont_var)

```



Splitting the data

```
x = df.drop(["RainTomorrow"], axis=1)
y = df["RainTomorrow"]

from sklearn.preprocessing import StandardScaler
scale=StandardScaler()
scale.fit(x)
StandardScaler()
X= scale.transform(x)
```

```

x.columns

Index(['Location', 'MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation',
      'Sunshine',
      'WindGustDir', 'WindGustSpeed', 'WindDir9am', 'WindDir3pm',
      'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am', 'Humidity3pm',
      'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm',
      'Temp9am',
      'Temp3pm', 'RainToday', 'Date_month', 'Date_day'],
      dtype='object')

X=pd.DataFrame(X,columns=x.columns)

X.head()

{"type": "dataframe", "variable_name": "X"}

y.head()

0    False
1    False
2    False
3    False
4    False
Name: RainTomorrow, dtype: bool

X_train, X_test, y_train, y_test = train_test_split(x,y, test_size
=0.2, random_state = 0)

```

Model Building using ML Models.

- RandomForestClassifier
- GaussianNB
- KNeighborsClassifier
- XGB Classifier

Random Forest Classifier

```

from sklearn.ensemble import RandomForestClassifier

ranfor= RandomForestClassifier()

ranfor.fit(X_train,y_train)

RandomForestClassifier()

ypred= ranfor.predict(X_test)

print(confusion_matrix(y_test,ypred))
print(accuracy_score(y_test,ypred))
print(classification_report(y_test,ypred))

```

```
[[21786  940]
 [ 3300 3066]]
0.8542554654200467
```


	precision	recall	f1-score	support
False	0.87	0.96	0.91	22726
True	0.77	0.48	0.59	6366
accuracy			0.85	29092
macro avg	0.82	0.72	0.75	29092
weighted avg	0.85	0.85	0.84	29092

```
from sklearn.metrics import RocCurveDisplay, roc_auc_score

# Plot ROC Curve
RocCurveDisplay.from_estimator(ranfor, X_test, y_test)

# Show the plot
plt.show()

# Compute AUC Score
print("AUC Score:", roc_auc_score(y_test, ypred))
```



AUC Score: 0.7201293979343144


Gaussian NB

```
from sklearn.naive_bayes import GaussianNB
gnb= GaussianNB()
gnb.fit(X_train,y_train)
GaussianNB()
ypred2= gnb.predict(X_test)
print(confusion_matrix(y_test,ypred2))
print(accuracy_score(y_test,ypred2))
print(classification_report(y_test,ypred2))
```

[[19844 2882]
 [2708 3658]]
0.8078509555891654

	precision	recall	f1-score	support
False	0.88	0.87	0.88	22726
True	0.56	0.57	0.57	6366
accuracy			0.81	29092
macro avg	0.72	0.72	0.72	29092
weighted avg	0.81	0.81	0.81	29092

```
from sklearn.metrics import RocCurveDisplay, roc_auc_score
# Plot ROC Curve
RocCurveDisplay.from_estimator(gnb, X_test, y_test)
# Show the plot
plt.show()
# Compute AUC Score
print("AUC Score:", roc_auc_score(y_test, ypred2))
```



AUC Score: 0.7239000206506067

K Nearest Neighbors

```
from sklearn.neighbors import KNeighborsClassifier
knn = KNeighborsClassifier(n_neighbors=3)
knn.fit(X_train,y_train)
KNeighborsClassifier(n_neighbors=3)
ypred3= knn.predict(X_test)
print(confusion_matrix(y_test,ypred3))
print(accuracy_score(y_test,ypred3))
print(classification_report(y_test,ypred3))
```

```
[[20939  1787]
 [ 3186 3180]]
0.8290595352674275
```

	precision	recall	f1-score	support
False	0.87	0.92	0.89	22726
True	0.64	0.50	0.56	6366


accuracy			0.83	29092
macro avg	0.75	0.71	0.73	29092
weighted avg	0.82	0.83	0.82	29092

```
from sklearn.metrics import RocCurveDisplay, roc_auc_score

# Plot ROC Curve
RocCurveDisplay.from_estimator(knn, X_test, y_test)

# Show the plot
plt.show()

# Compute AUC Score
print("AUC Score:", roc_auc_score(y_test, ypred3))
```



AUC Score: 0.7104481715255037

XGB Classifier

```
from xgboost import XGBClassifier

xgb= XGBClassifier()
```



```

xgb.fit(X_train,y_train)

XGBClassifier(base_score=None, booster=None, callbacks=None,
              colsample_bylevel=None, colsample_bynode=None,
              colsample_bytree=None, device=None,
early_stopping_rounds=None,
              enable_categorical=False, eval_metric=None,
feature_types=None,
              gamma=None, grow_policy=None, importance_type=None,
              interaction_constraints=None, learning_rate=None,
max_bin=None,
              max_cat_threshold=None, max_cat_to_onehot=None,
              max_delta_step=None, max_depth=None, max_leaves=None,
              min_child_weight=None, missing=nan,
monotone_constraints=None,
              multi_strategy=None, n_estimators=None, n_jobs=None,
              num_parallel_tree=None, random_state=None, ...)

ypred4= xgb.predict(X_test)

print(confusion_matrix(y_test,ypred4))
print(accuracy_score(y_test,ypred4))
print(classification_report(y_test,ypred4))

[[21508  1218]
 [ 2896  3470]]
0.8585865530042623

```

	precision	recall	f1-score	support
False	0.88	0.95	0.91	22726
True	0.74	0.55	0.63	6366
accuracy			0.86	29092
macro avg	0.81	0.75	0.77	29092
weighted avg	0.85	0.86	0.85	29092

```


from sklearn.metrics import RocCurveDisplay, roc_auc_score

# Plot ROC Curve
RocCurveDisplay.from_estimator(xgb, X_test, y_test)

# Show the plot
plt.show()

# Compute AUC Score
print("AUC Score:", roc_auc_score(y_test, ypred4))

```



AUC Score: 0.7457441267355018

save the best performing model i.e. XGB Classifier model in our pickle file

```
import pickle

# Save the model
with open("rain_XGBnew_model.pkl", "wb") as file:
    pickle.dump(xgb, file)

# Load the model
with open("rain_XGBnew_model.pkl", "rb") as file:
    model = pickle.load(file)
```